

WAIT TIME MONITORING FOR CAMPUS EATERIES

USING COMPUTER VISION

PRESENTED BY : TEAM 011

Ayush Trivedi
Dheeraj Pamnani
Dominic Darrah
Riya Agarwal
Sravani Bolla



1. INTRODUCTION - PROBLEM IDENTIFICATION

- Unpredictable campus eatery wait times causing frustration..
- At ASU, students often experience long, unpredictable wait times at popular campus eateries like Starbucks, Chick-fil-A, and Pitchforks.
- These delays cause frustration, missed class time, and inefficient use of short breaks.
- There is currently **no digital system** providing real-time visibility into queue lengths.



Why We Chose This Problem:

- It directly impacts a large part of the ASU community (students, staff, visitors).
- Solving it improves daily productivity, well-being, and campus operations.
- The problem is **feasible** to address using Computer Vision (CV) while maintaining privacy.
- There is a clear gap — no existing wait-time tracking solution is implemented at ASU.

Research Conducted:

- Observed dining queues, interviewed students, reviewed ASU's tech (no live tracking), and benchmarked external real-time crowd monitoring systems using computer vision.

1. INTRODUCTION - PROBLEM IDENTIFICATION

- Unpredictable campus eatery wait times causing frustration.

How It Is Being Addressed Currently:

- Students check lines manually by walking to locations.
- They sometimes rely on word-of-mouth from friends or text updates.
- Vendors do not provide real-time queue information or predictive tools.



Issues with the Current State:

- **Inefficiency:** Students waste time traveling to crowded locations and then leaving.
- **Frustration:** Lack of information causes anxiety, stress, and poor time management.
- **Lost Revenue:** Vendors miss sales opportunities from students abandoning long lines.
- **Operational Blind Spot:** Dining management cannot optimize staffing or manage peak traffic without real-time data.
- **Predictive Capability:** ASU cannot predict or balance traffic across vendors, leading to localized overcrowding during class breaks.

2. PROJECT DEFINITION

Scope of the Solution:

- Monitor real-time queue lengths at high-traffic ASU eateries (e.g. Starbucks, Chick-fil-A, Pitchforks).
- Use overhead or wall-mounted cameras to capture queue footage.
- Analyze footage using a Computer Vision (YOLO v8) model to identify and count people to estimate wait times.
- Display live wait times via the ASU Mobile App and digital signage around campus.

Value of the Solution:

- **Utility:** Allows students to plan meals efficiently, reducing wasted trips and time.
- **Impact:** Enhances student experience, improves vendor operations, and fosters a tech-forward campus image.
- **Cost-Benefit:** Upfront investment in cameras and server infrastructure is offset by increased dining revenue, reduced congestion, and better campus resource management.

Key Stakeholders and Beneficiaries

Students and Visitors

Dining Services and Vendors

ASU Administration

App Developers

2. PROJECT DEFINITION

- Real-time queue tracking using cameras and computer vision.

Success Metrics:

- **Student Adoption:** Majority of students reporting satisfaction with queue visibility tools.
- **Operational Impact:** Significant reduction in average wait times during peak hours.
- **Vendor Sales:** Increase in transaction numbers during historically busy periods.

References for case studies are provided in slide 19
(Reference slide)

How Our Solution Is Different:

- **Privacy-Focused:** Our system avoids identity recognition, using only anonymous body detection.
- **University Context:** Adapted to dynamic, fast-paced campus environments with highly variable peak times.
- **Cost-Efficient:** Uses lightweight, open-source CV models like YOLO v8 instead of expensive proprietary sensor network

Case Studies:



Disney Parks

Use CV and sensors to monitor ride wait times and direct crowd flow.



Airports

Real-time TSA security checkpoint wait times using CV and Bluetooth tracking.



Theme Parks & Events

Crowd management through camera analytics and mobile app updates.

3. DATA ACQUISITION/ PREPARATION

- Collect and prepare cafeteria queue images for model training.

Data Sources:

- Images of cafeteria queues (Starbucks, Chick-fil-A, Pitchforks).
- Focused on varied crowd sizes, lighting conditions, and times of day.

Data Collection:

- Manually captured images on campus (with privacy protection).
- Selected frames where queues were clearly visible.

Cleaning and Preprocessing:

Removed blurry or irrelevant frames.

Resize all images to 640x640 pixels for model input.

Annotated bounding boxes around people as ground-truth labels.



3. DATA ACQUISITION/ PREPARATION

- Collect and prepare cafeteria queue images for model training.

Challenges During Data Preparation:

Lighting Variations

- Natural lighting at different eateries changed drastically (sunlight vs indoor lights).



Crowded Scenes

- Overlapping people made bounding box labeling harder.



Privacy Concerns

- Ensured all images focused on body outlines, not faces, to protect individual identity..



4. FEATURE/LABEL ENGINEERING

- Create features and labels by annotating people in queue images.

Preprocessing Steps:

Image Resizing

All images resized to 640x640 pixels for YOLOv8 input.



Frame Cleaning

Dropped blurry, empty, or irrelevant frames to maintain high data quality.



Normalization

Scaled pixel values to [0,1] range for faster and more stable model training.

Features and Labels for the Model:

Features:

- Full-frame color images capturing cafeteria queues.

Labels:

- Manually annotated bounding boxes around every detected person based on shops (single class: "person").
- Bounding boxes used to train the object detection model for people-counting, not identity recognition.

5. MODEL EVALUATION

- Tested model accuracy and gathered user feedback.

PRE - DEPLOYMENT

- **Model Accuracy Testing:**
 - Use a hold-out test set of cafeteria queue images.
- **Manual Verification:**
 - Compare predicted people counts against manually counted ground truth.
 - Validate on different lighting conditions (indoor, outdoor) and crowd densities.
- **User Pilot Testing:**
 - Soft launch with a few eateries and collect student feedback on estimated wait times

POST - DEPLOYMENT

- **Continuous Accuracy Monitoring:**
 - Randomly sample camera feeds weekly for manual spot checks.
 - Calculate model Precision, Recall periodically to detect model drift.
- **User Feedback Surveys:**
 - Regular surveys inside the ASU app asking about perceived wait-time accuracy.
- **Operational Metrics:**
 - Monitor changes in average wait times and vendor transaction counts after system implementation.
- **Retraining Plan:**
 - Schedule retraining every semester using newly collected cafeteria queue data.

6. DEPLOYMENT

- Installed cameras, integrated system, and launched app updates.

Estimated People, Data, Systems, and Computational Resources/Costs:

People

- 1–2 CV Engineers (model training, validation).
- 1 Backend Developer (API, integration with ASU app).
- 1 IT Technician (camera installation/maintenance).

Data

- Initial training set: ~3,000 annotated images of cafeteria queues.
- Ongoing data collection every semester (~1,000 new images/semester).

Systems

- Overhead cameras (one per key location): ~\$200–\$400 each.
- Central server (cloud-hosted GPU instance or ASU campus server).
- Storage for images/videos (~1 TB/year).

Computational Resources

- Training Phase: Cloud GPU (e.g., AWS EC2 p3 instance) for model development.
- Deployment Phase: Lightweight edge computing devices or backend inference server.

Estimated Costs:

- Initial setup: ~\$8,000–\$12,000 (cameras + server + manpower).
- Annual maintenance: ~\$3,000 (camera upkeep, server costs, retraining labor).

6. DEPLOYMENT

- Installed cameras, integrated system, and launched app updates.

Estimated Workflow Updates/Costs (People, Process, Technology):

People

- Dining staff minor training: how to report camera issues (no operational burden).
- IT support for periodic system checks.

Process Changes

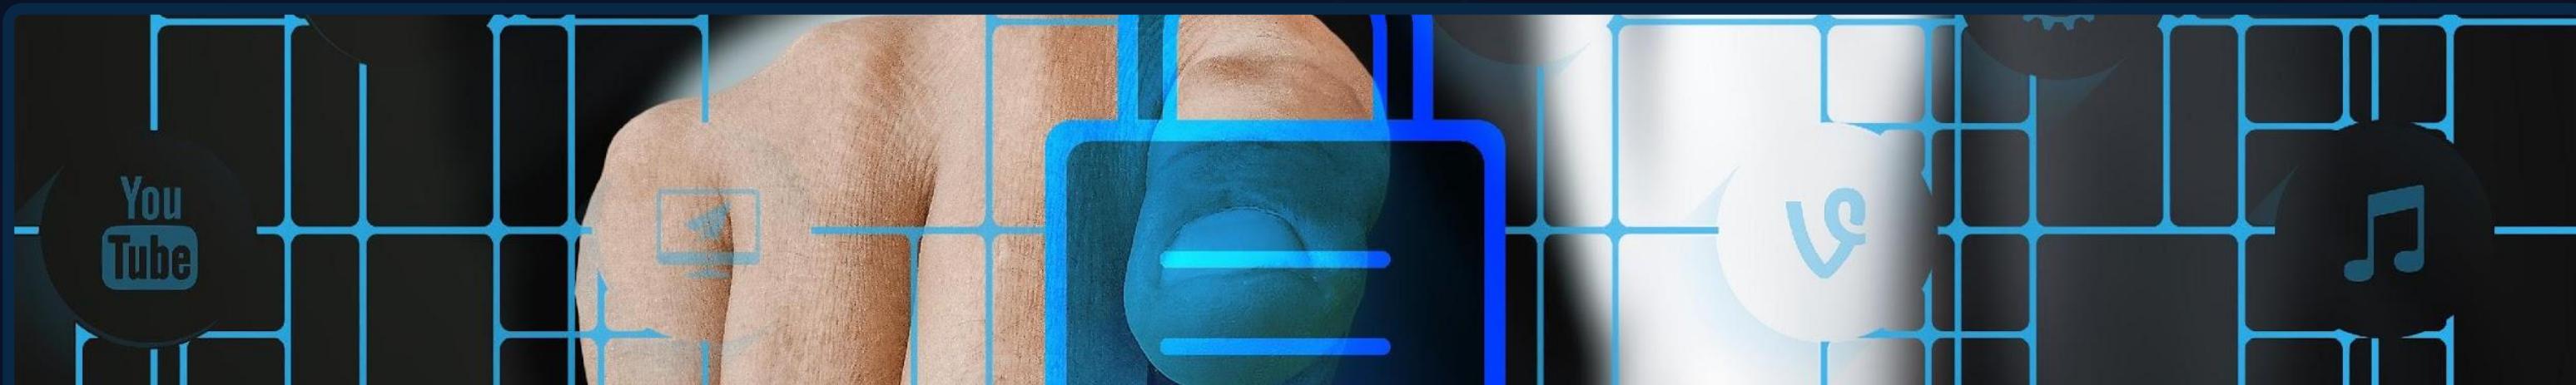
- Update ASU Mobile App to show real-time queue wait times.
- Display queue lengths on digital signage screens inside Student Union buildings.

Technology Changes

- Camera system installation and network connectivity setup.
- Integration of CV model inference output with backend servers and ASU APIs.

Overall Cost-Benefit:

- Initial investment recouped through improved vendor sales, better space utilization, and increased student satisfaction



7. MONITORING/MAINTENANCE

Monitoring and Updating:

- Continuous system checks and model retraining.
- Weekly manual spot checks of camera feeds.
- Monthly performance reviews (Precision, Recall metrics).
- Semester-based model retraining with new cafeteria data.
- Continuous user feedback collection via ASU App surveys.

Unintended Incentives/Consequences (Mitigation)

- **Sudden Crowding:** Students rush vendors showing "short wait" → Smooth updates, randomized refresh times.
- **Vendor Manipulation:** Vendors may alter service speed → Cross-validate queue estimates with transaction logs.

Privacy and Security (Mitigation)

- **Risk:** Potential misuse of visual data.
- **Mitigation:** No raw video storage; real-time detection only; encrypted data transmission; no facial recognition.

Lessons Learned

- Scope narrowed from "wait time prediction" to "real-time queue monitoring" for feasibility.
- Discovered lighting and crowding issues, requiring advanced data augmentation and model tuning.

Scalability

- Expand to all ASU campuses and venues.
- Future add-ons: Wait time prediction, event crowd monitoring (libraries, gyms, etc.).

COMPUTER VISION MODEL

CV Model Description

- **Model Used:** YOLO v8 (Ultralytics, Open-Source).
- **Model Design:**
 - a. One-stage object detector.
 - b. Predicts bounding boxes and confidence scores directly from input images.
 - c. Optimized for fast inference (real-time performance) and small objects (crowded scenes).



Which Part of the Solution Uses CV

- CV is used to **analyze camera footage** from cafeteria locations to **detect and count the number of people** in a queue in real-time.
- Output: Live queue length estimation → fed into backend → shown in ASU app and digital signage.

Why CV Is Required

- Manual counting (by staff or students) is unreliable, slow, and not scalable.
- Sensor-based alternatives (entry counters) can't distinguish people **waiting** vs **walking through**.
- Only Computer Vision can provide **accurate, anonymous, real-time queue analysis**.
- Without CV, real-time detection, automation, and dynamic updates would be impossible.

COMPUTER VISION MODEL

Model Training Details

- **Data:** ~100 annotated images of cafeteria queues (custom and open-source).
- **Preprocessing:** Image resizing and cropping for robustness.
- **Model:** YOLO v8 (pre-trained on COCO, fine-tuned on our queue dataset).

External Validation Plan (Pre and Post Deployment)

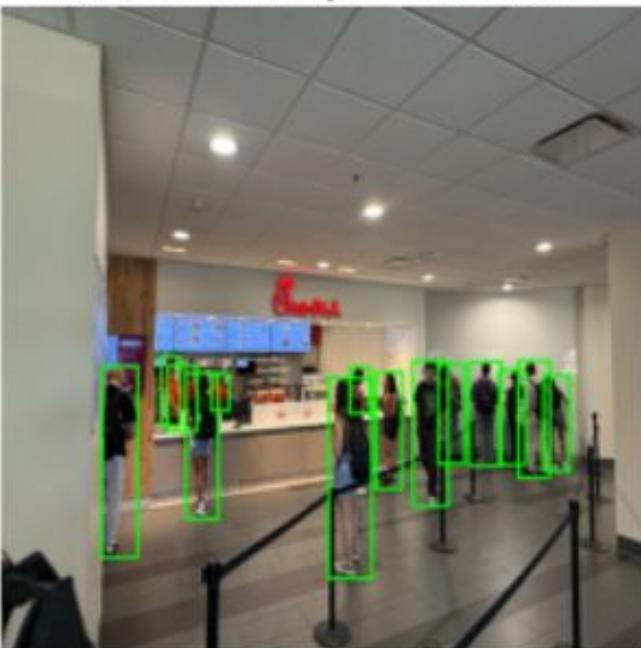
- Pilot launch at selected eateries; real-world people counts compared against manual counts.
- Ongoing weekly random manual audits post-deployment.

Outcome-Action Pairings

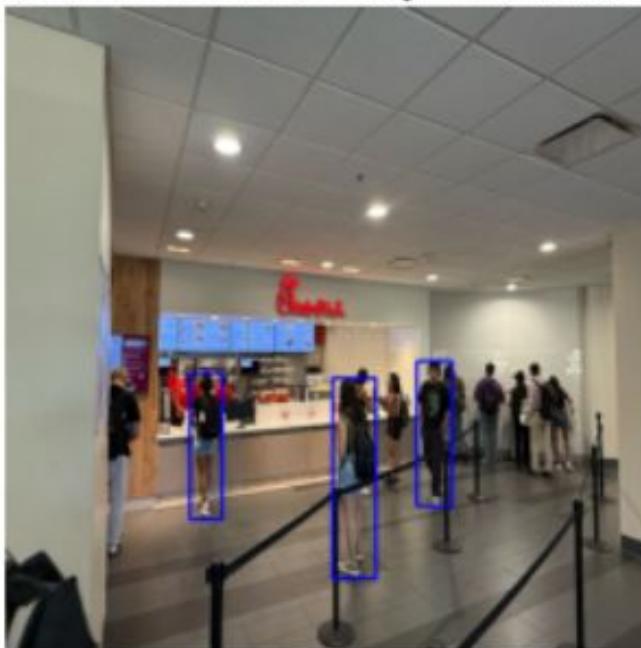
- **TP (True Positive):** Correct people detection → Accurate wait time displayed.
- **FP (False Positive):** Wrong detection → Slightly inflated wait time (acceptable).
- **FN (False Negative):** Missed detection → Slightly underestimated wait time (monitored).
- **TN (True Negative):** Correct no detection → No queue shown.

WORKING DEMO

All Detected People - Chick fila a



Filtered Customers Only - Chick fil a



Average wait time per customer: 29.19 seconds
Standard deviation: 5.88 seconds

	shop	customer_count	predicted_wait_time(seconds)	predicted_wait_time(minutes)
0	Chick fila a	24	700.50	11.68
1	Choolah	1	29.19	0.49
2	Gigis	5	145.94	2.43
3	POD Market	16	467.00	7.78
4	Pei-Wei	6	175.12	2.92
5	QDoba	25	729.69	12.16
6	Starbucks	10	291.88	4.86
7	burger king	4	116.75	1.95
8	einstein bros bagels	5	145.94	2.43
9	subway	0	0.00	0.00

- Proof of Concept Demo
 - Walk-through of Computer Vision Model
 - Accurately identify customers in queue to order
 - Predicting wait time for customer
 - **Wait time** refers to the duration between a customer joining the eatery line and placing their order.
 - Initially, the average wait time was assumed to be 30 seconds with a standard deviation of 10 seconds. After analyzing annotated images based on customer count, the observed average wait time per customer was 29.19 seconds with a standard deviation of 5.88 seconds.



CONCLUSION



1.

SOLUTION OVERVIEW

- Built a real-time Computer Vision system to detect and count people in cafeteria queues at ASU.
- Deployed YOLOv8 model on camera feeds to estimate live wait times, integrated with ASU app and digital signage.
- Focused on anonymous, privacy-preserving detection without storing facial or identity data.

2.

KEY FINDINGS

- Real-time queue updates improve student satisfaction and dining vendor efficiency.
- Computer Vision approach is scalable and operationally practical for campus environments.
- Shops like Chick-fil-A and Qdoba have the highest customer counts and predicted wait times compared to others.
- Smaller shops like Choolah, Burger King, and Einstein Bros Bagels have fewer customers and much shorter predicted wait times.

CONCLUSION

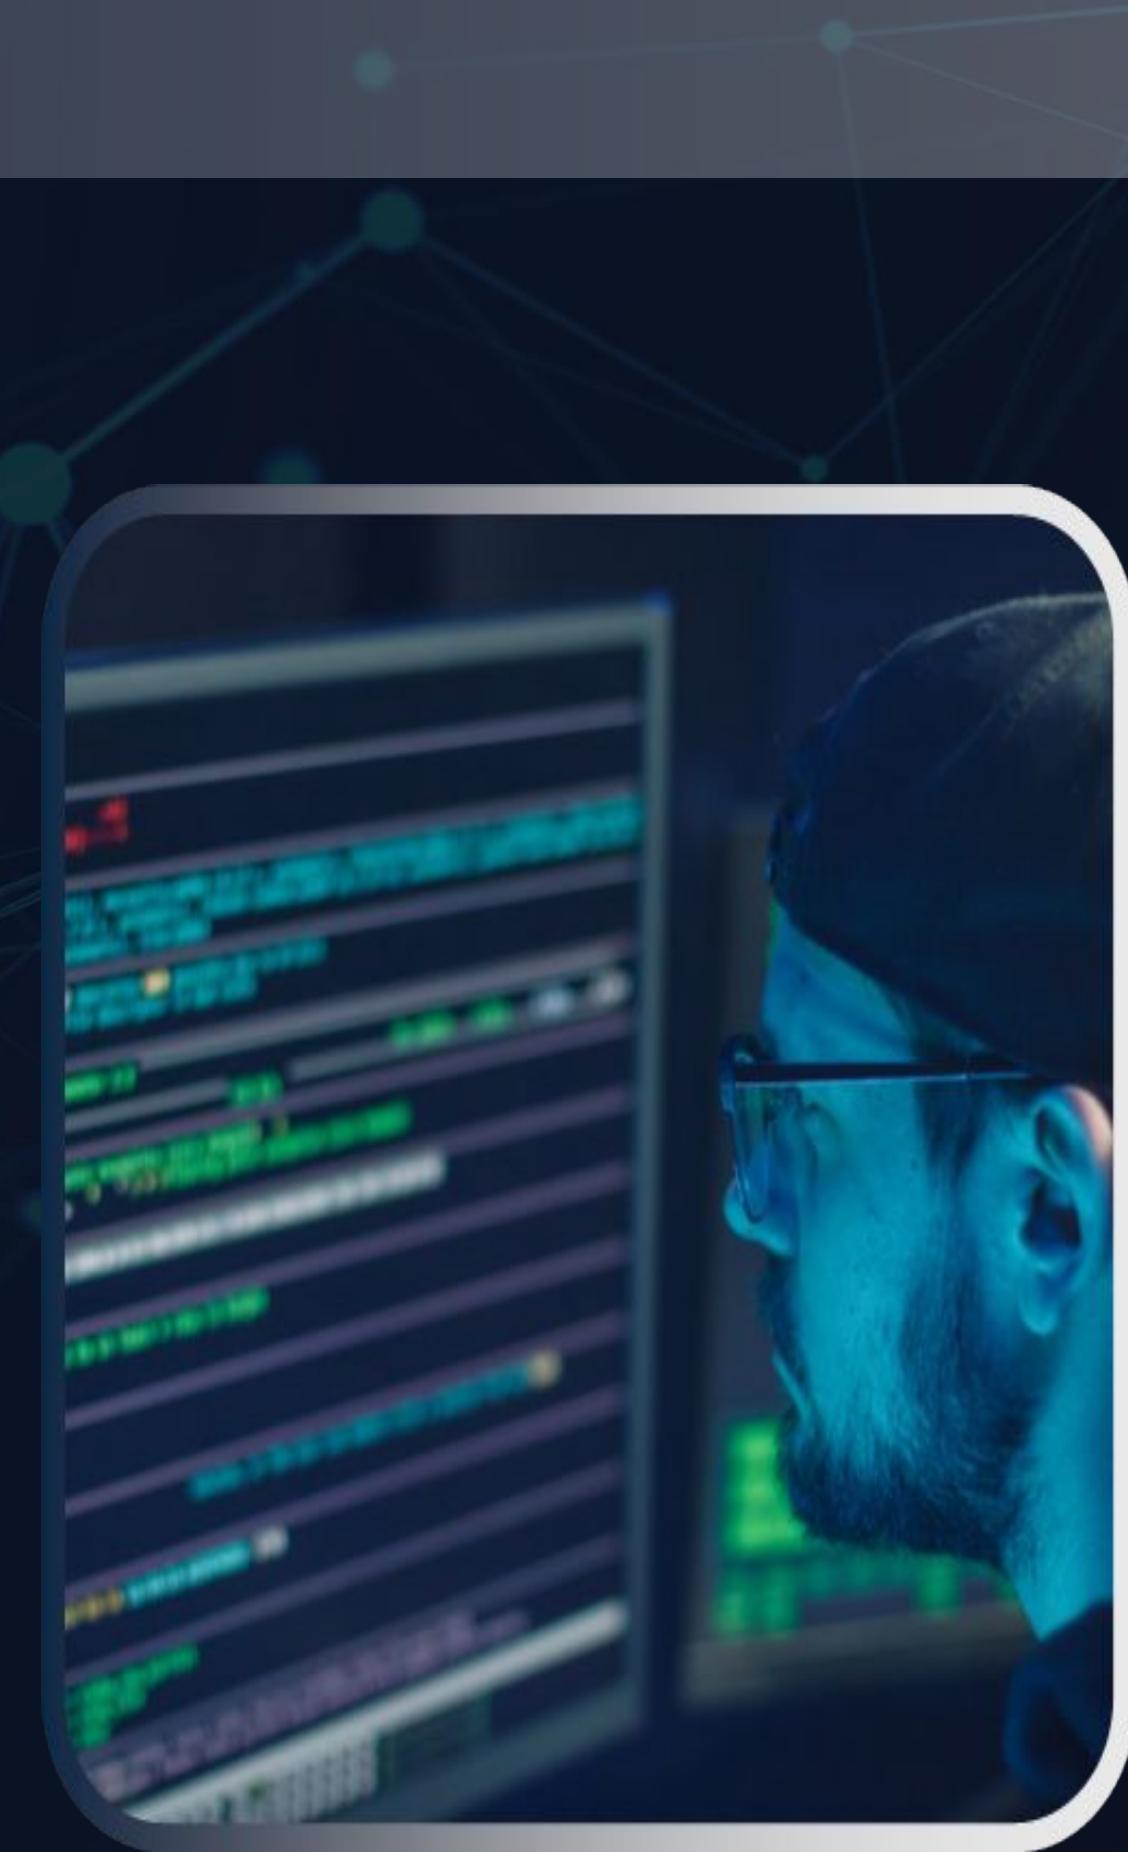
continued

3. LIMITATIONS

- Reduced accuracy in dense, poorly lit environments; occasional misses due to occlusion or non-ideal poses.
- Static cameras miss dynamic queue movements without advanced tracking.
- Struggles in dense crowds; plan to use instance segmentation models like Mask R-CNN.
- Current pre-trained YOLOv8 models miss occasional customers standing in queues, due to non-ideal poses or occlusion.

4. FUTURE WORK

- Integrate predictive analytics to estimate future wait and service times based on historical patterns and transaction rates.
- Expand the system to libraries, gym facilities, event spaces, and other areas.
- Upgrade detection models by using instance segmentation models like Mask R-CNN for better performance in crowded scenes.
- Develop a continuous learning system with live feedback and auto-retraining for improved accuracy.
- Fine-tune models specifically for university eatery queues and extend to handle dynamic moving queues.



Tabulated Task Summary

Team Member	CV Model / Research	Presentation Content Creation	Final Presentation Delivery
Ayush Trivedi	YOLO v8 model implementation, tuning	Slides 15-17	Conducted live demo and conclusions
Dheeraj Pamnani	Model validation - outcome scenarios; wait time calculation	Slides 12-14	Walked through monitoring plan and Computer Vision model architecture
Dominic Darrah	Data collection, cleaning, annotation strategy, privacy considerations	Slides 6-8	Presented data acquisition & preprocessing steps
Riya Agarwal	Problem scope identification, value and impact analysis	Slides 2-5	Opened the presentation and discussed impact & value add
Sravani Bolla	Model validation, pre/post deployment metrics; proposed resource plan	Slides 9-11	Explained validation process & deployment pipeline

Each member played an active role in the early-stage discussions, from identifying the core problem to ideating the project solution.

References

- <https://www.ultralytics.com/blog/revolutionizing-queue-management-with-ultralytics-yolov8-and-openvino>
- <https://github.com/ultralytics/ultralytics>
- [Blackwells Capital Outlines Plans for AI at Disney Theme Parks \(Crowd Management, Predictive Maintenance, Real-Time Ticket Price Changes\)](#)
- [CrowdVision - Automated passenger tracking using video analytics](#)
- [YOLOv11 Real-time Queue Management | Ultralytics](#)
- AI LLM to assist in syntax

THANK YOU!

